

Guidebook: Building a CI/CD Pipeline for a Python App on AWS EKS

This guide details the process of creating an automated pipeline that takes a Python application from a Git commit to a live deployment on a Kubernetes cluster.

Phase 1: Local Setup & Source Control

1. Create Project Directory & Files:

- Open a terminal and create a project folder:

```
mkdir cicd-project
cd cicd-project
```

- Create `main.py` with your FastAPI application:

```
from fastapi import FastAPI
app = FastAPI()
@app.get("/")
def read_root():
    return {"message": "Hello, World! CI/CD Pipeline is running."}
```

- Create `requirements.txt` for dependencies:

```
fastapi
uvicorn[standard]
```

2. Set Up GitHub Repository:

- Go to GitHub and create a new, empty repository (e.g., `cicd-app-repo`).
- In your local project folder, initialize Git and push your code. Replace the URL with your own.

```
git init -b main
git add .
git commit -m "Initial commit"
git remote add origin https://github.com/YOUR_USERNAME/cicd-app-repo.git
git push -u origin main
```

Phase 2: Containerization with Docker

1. Create Docker-related Files:

- In your project root, create a `Dockerfile`:

```
# Start with an official lightweight Python image
FROM python:3.9-slim
```

```
# Set the working directory inside the container
WORKDIR /app
```

```
# Copy the dependencies file first to leverage Docker's layer caching
COPY requirements.txt .
```

```
# Install the dependencies
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
# Copy the rest of the application code into the container  
COPY . .
```

```
# Command to run the application when the container starts
```

```
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

- Create a `.dockerignore` file to keep the image small:

```
__pycache__/  
*.pyc  
.git  
.vscode  
Venv
```

2. Test Docker Image Locally:

- Build the image:
`docker build -t cicd-app .`
- Run the container:
`docker run -p 8000:8000 cicd-app`
- Verify it's working by opening `http://localhost:8000` in your browser.

3. Commit Docker Files to GitHub:

```
git add Dockerfile .dockerignore  
git commit -m "feat: Add Docker containerization"  
git push origin main
```

Phase 3: Initial AWS Setup (IAM & ECR)

1. Create an IAM User for CLI Access:

- In the AWS Console, go to **IAM -> Users -> Create user**.
- Name the user (e.g., `cli-user`) and attach the `AdministratorAccess` policy for this project. **Note:** In a real production environment, you would grant more restrictive permissions.
- Create the user and go to the **Security credentials** tab. Create an **access key** and save the `Access Key ID` and `Secret Access Key` securely.

2. Create ECR Repository:

- Go to the **Amazon ECR** service.
- Create a **private** repository named `cicd-app`.

Phase 4: Building the CI Pipeline

1. Create the Build Specification (`buildspec.yml`):

- In your local project, create `buildspec.yml`:

version: 0.2

phases:

pre_build:

commands:

- echo Logging in to Amazon ECR...
- aws ecr get-login-password --region \$AWS_DEFAULT_REGION | docker login

--username AWS --password-stdin

\$AWS_ACCOUNT_ID.dkr.ecr.\$AWS_DEFAULT_REGION.amazonaws.com

build:

commands:

- echo Build started on `date`
- echo Building the Docker image...
- docker build -t \$IMAGE_REPO_NAME:\$IMAGE_TAG .
- docker tag \$IMAGE_REPO_NAME:\$IMAGE_TAG

\$AWS_ACCOUNT_ID.dkr.ecr.\$AWS_DEFAULT_REGION.amazonaws.com/\$IMAGE_REPO_NAME:\$IMAGE_TAG

post_build:

commands:

- echo Build completed on `date`
- echo Pushing the Docker image to ECR...
- docker push

\$AWS_ACCOUNT_ID.dkr.ecr.\$AWS_DEFAULT_REGION.amazonaws.com/\$IMAGE_REPO_NAME:\$IMAGE_TAG

- Commit and push this file to GitHub.

git add buildspec.yml

git commit -m "feat: Add buildspec for CodeBuild"

git push origin main

2. Create the CodeBuild Project:

- In the AWS Console, go to **CodeBuild** -> **Create build project**.
- **Project name:** cicd-app-build.
- **Source:** Connect to your GitHub repository and select your cicd-app-repo. Enable the webhook to rebuild on every code push.
- **Environment:** Use a **Managed image** (Amazon Linux 2, Standard). **Crucially, check the "Privileged" box.**
- **Environment variables:** Add the following:
 - AWS_ACCOUNT_ID: *Your 12-digit AWS Account ID*
 - AWS_DEFAULT_REGION: *Your AWS region code (e.g., ap-south-1)*
 - IMAGE_REPO_NAME: cicd-app
 - IMAGE_TAG: latest
- Create the build project.

3. Grant ECR Permissions:

- After the project is created, a new IAM role will be generated. Go to the **IAM** console, find this role (codebuild-cicd-app-build-service-role), and attach the

AmazonEC2ContainerRegistryPowerUser policy.

4. Create the CodePipeline:

- Go to **CodePipeline** -> **Create pipeline**.
- **Name:** cicd-app-pipeline.
- **Source stage:** Select **GitHub (Version 2)** and choose your repository and main branch.
- **Build stage:** Select **AWS CodeBuild** and choose the cicd-app-build project you just created.
- **Deploy stage:** **Skip** this stage for now.
- Create the pipeline. It will run automatically and push an image to your ECR repository.

Phase 5: Kubernetes Cluster Setup (EKS)

1. Create the EKS Cluster:

- Go to **Amazon EKS** -> **Create cluster**.
- **Name:** cicd-cluster.
- Follow the wizard to create the necessary **Cluster IAM role**.
- In the **Compute** section, create a **Node group**. Let the wizard create the **Node IAM role**. Use t3.small instances and set the desired size to 2.
- Create the cluster. This will take 10-20 minutes.

2. Tag Public Subnets for Load Balancer:

- While the cluster is creating, go to the **VPC** console.
- Find the VPC created for your cluster. Go to **Subnets**.
- Identify your **public subnets** (those with a route to an Internet Gateway igw-).
- For each public subnet, go to the **Tags** tab and add the following tag:
 - **Key:** kubernetes.io/role/elb
 - **Value:** 1

Phase 6: Connecting kubectl to EKS

1. Configure Local CLI:

- Install the AWS CLI and kubectl on your machine.
- Configure your local AWS CLI with the credentials of the IAM user you created in Phase 3.

```
Bash
```

```
aws configure
```

2. Set Up Cluster Access (aws-auth):

- This step must be done from **AWS CloudShell**, which has initial admin access.
- Open CloudShell and create a file named aws-auth-cm.yaml with the complete, correct configuration:

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
  mapRoles: |
    - rolearn: YOUR_NODE_IAM_ROLE_ARN
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
  mapUsers: |
    - userarn: YOUR_CLI_USER_ARN
      username: cli-user
      groups:
        - system:masters

```

Replace YOUR_NODE_IAM_ROLE_ARN and YOUR_CLI_USER_ARN with the actual ARNs.

- Apply this configuration in CloudShell:

```
kubectl apply -f aws-auth-cm.yaml
```

3. Connect Local kubectl:

- Back on your local machine, run the command to update your kubeconfig file:

```
aws eks update-kubeconfig --region YOUR_REGION --name cisd-cluster
```

- Test the connection:

```
kubectl get nodes
```

You should see your two nodes listed.

Phase 7: Deploying the Application

1. Create Manifest Files:

- In your local project folder, create deployment.yaml:

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: cisd-app-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: cisd-app
  template:
    metadata:
      labels:
        app: cisd-app

```

```
spec:
  containers:
    - name: cicd-app
      image: YOUR_AWS_ACCOUNT_ID.dkr.ecr.YOUR_REGION.amazonaws.com/cicd-app:latest
      ports:
        - containerPort: 8000
```

- Create `service.yaml`, including the annotation to ensure an internet-facing load balancer:

```
apiVersion: v1
kind: Service
metadata:
  name: cicd-app-service
  annotations:
    service.beta.kubernetes.io/aws-load-balancer-scheme: internet-facing
spec:
  selector:
    app: cicd-app
  type: LoadBalancer
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8000
```

Remember to replace the placeholders in `deployment.yaml`.

2. Apply Manifests:

- From your local terminal, apply the manifests:

```
kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
```

Phase 8: Verification

1. Check Pods and Service:

```
kubectl get pods
kubectl get service cicd-app-service
```

2. Wait a few minutes for the `EXTERNAL-IP` of the service to be assigned.
3. Copy the DNS name and paste it into your browser. You should see your application running.

Phase 9: Project Cleanup

To avoid costs, delete all the resources you created, starting with the resources in Kubernetes and then moving to AWS.

1. **Delete K8s LoadBalancer & Deployment:** `kubectl delete service cicd-app-service` and `kubectl delete deployment cicd-app-deployment`.
2. **Delete EKS Cluster:** In the AWS Console.
3. **Delete ECR Repository:** In the AWS Console.

4. **Delete CodePipeline & CodeBuild Project:** In the AWS Console.
5. **Delete S3 bucket** created by CodePipeline.
6. **Delete IAM Roles** created during the setup.